

```
{"nbformat":4,"nbformat_minor":0,"metadata":{"colab":{"provenance":[]}},"kernelspec": {"name":"python3","display_name":"Python 3","language_info":{"name":"python"}}, "cells": [{"cell_type":"code","metadata":{"id":"mCGU3AhSlp2W"},"executionInfo": {"status":"ok","timestamp":1769770079790,"user_tz":180,"elapsed":24864,"user": {"displayName":"Luciano de Andrade Barbosa","userId":"01495378349148092861"}}, "source": ["# mlp for multiclass classification\n", "from numpy import argmax\n", "from pandas import read_csv\n", "from sklearn.model_selection import train_test_split\n", "from sklearn.preprocessing import LabelEncoder\n", "from tensorflow.keras import Sequential\n", "from tensorflow.keras.layers import Dense"], "execution_count":1,"outputs":[]}, {"cell_type":"code","metadata":{"colab": {"base_uri":"https://localhost:8080/"},"id":"71tLw6DVmeo3","outputId":"cbf2ba1a-3a69-4a4e-aa11-337de33dab39"}, "executionInfo": {"status":"ok","timestamp":1769770092129,"user_tz":180,"elapsed":169,"user": {"displayName":"Luciano de Andrade Barbosa","userId":"01495378349148092861"}}, "source": ["# load the dataset\n", "path =\n", "https://raw.githubusercontent.com/jbrownlee/Datasets/master/iris.csv\n", "df = read_csv(path,\n", "header=None)\n", "# split into input and output columns\n", "X, y = df.values[:, :-1], df.values[:,\n", "-1]\n", "# ensure all data are floating point values\n", "X = X.astype('float32')\n", "# encode strings\n", "to integer\n", "y = LabelEncoder().fit_transform(y)\n", "# split into train and test\n", "datasets\n", "X_train, X_test, y_train, y_test = train_test_split(X, y,\n", "test_size=0.33)\n", "print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)\n", "#\n", "determine the number of input features\n", "n_features =\n", "X_train.shape[1]"], "execution_count":2,"outputs": [{"output_type":"stream","name":"stdout","text": ["(100, 4) (50, 4) (100,) (50,)"]}], {"cell_type":"code","metadata": {"id":"KEzTT8ogmmoG","colab":{"base_uri":"https://localhost:8080/"}}, "executionInfo": {"status":"ok","timestamp":1769770210248,"user_tz":180,"elapsed":350,"user": {"displayName":"Luciano de Andrade\nBarbosa","userId":"01495378349148092861"}}, "source": ["# define model\n", "model = Sequential()\n", "model.add(Dense(10,\n", "activation='relu', kernel_initializer='he_normal', input_shape=\n", "(n_features,)))\n", "model.add(Dense(8, activation='relu',\n", "kernel_initializer='he_normal'))\n", "model.add(Dense(3,\n", "activation='softmax'))"], "execution_count":3,"outputs": [{"output_type":"stream","name":"stderr","text": ["/usr/local/lib/python3.12/dist-\n", "packages/keras/src/layers/core/dense.py:93: UserWarning: Do not pass an\n", "'input_shape' / 'input_dim' argument to a layer. When using Sequential models, prefer using an\n", "'Input(shape)' object as the first layer in the model instead.\n", "\n", "super().__init__(activity_regularizer=activity_regularizer, **kwargs)"]}], {"cell_type":"code","metadata":{"colab": {"base_uri":"https://localhost:8080/"},"id":"haGlv1Nzmorl","outputId":"e827eccc-e3da-436e-b54e-2bc1f5d63d4c"}, "executionInfo": {"status":"ok","timestamp":1769770248708,"user_tz":180,"elapsed":9830,"user": {"displayName":"Luciano de Andrade Barbosa","userId":"01495378349148092861"}}, "source": ["# compile the model\n", "model.compile(optimizer='adam',\n", "loss='sparse_categorical_crossentropy', metrics=['accuracy'])\n", "# fit the\n", "model\n", "model.fit(X_train, y_train, epochs=150, batch_size=32,\n", "verbose=0)"], "execution_count":4,"outputs": [{"output_type":"execute_result","data": {"text/plain": [""]}], "metadata":{}, "execution_count":4}], {"cell_type":"code","metadata":{"colab": {"base_uri":"https://localhost:8080/"},"id":"K0oDNi36m0ns","outputId":"25d1acf2-2dda-4a66-ad74-f8f60f2fdc3c"}, "executionInfo": {"status":"ok","timestamp":1769770318966,"user_tz":180,"elapsed":210,"user": {"displayName":"Luciano de Andrade Barbosa","userId":"01495378349148092861"}}, "source": ["# evaluate the model\n", "loss, acc = model.evaluate(X_test, y_test, verbose=0)\n", "print('Test\n", "Accuracy: %.3f % acc)\n", "# make a prediction\n", "row = [5.1,3.5,1.4,0.2]\n", "import numpy as\n", "np\n", "yhat = model.predict(np.array([row]))\n", "print('Predicted: %s (class=%d) % (yhat,\n", "argmax(yhat)))"], "execution_count":6,"outputs": [{"output_type":"stream","name":"stdout","text": [""]}]}
```

```
["Test Accuracy: 0.900\n", "\u001b[1m1/1\u001b[0m\n\u001b[32m\u001b[0m\u001b[37m\u001b[0m\n\u001b[1m0s\u001b[0m 70ms/step\n", "Predicted: [[0.9263848 0.07129735 0.00231787]]\n(class=0)\n"]}, {"cell_type": "code", "metadata": {"colab": {"base_uri": "https://localhost:8080/"}, "id": "3DR7JYe4xDrC", "outputId": "99126356-7c00-4a9f-eb91-765ac8710fe4"}, "source": ["# mlp for regression\n", "from numpy import sqrt\n", "from pandas import read_csv\n", "from sklearn.model_selection import train_test_split\n", "from tensorflow.keras import Sequential\n", "from tensorflow.keras.layers import Dense\n", "# load the dataset\n", "path =\n", "https://raw.githubusercontent.com/jbrownlee/Datasets/master/housing.csv\n", "df =\n", "read_csv(path, header=None)\n", "# split into input and output columns\n", "X, y = df.values[:, :-1],\n", "df.values[:, -1]\n", "# split into train and test datasets\n", "X_train, X_test, y_train, y_test =\n", "train_test_split(X, y, test_size=0.33)\n", "print(X_train.shape, X_test.shape, y_train.shape,\n", "y_test.shape)\n", "# determine the number of input features\n", "n_features =\n", "X_train.shape[1]\n", "# define model\n", "model = Sequential()\n", "model.add(Dense(10,\n", "activation='relu', kernel_initializer='he_normal', input_shape=\n", "(n_features,)))\n", "model.add(Dense(8, activation='relu',\n", "kernel_initializer='he_normal'))\n", "# compile the\n", "model.compile(optimizer='adam', loss='mse')\n", "# fit the model\n", "model.fit(X_train,\n", "y_train, epochs=150, batch_size=32, verbose=0)\n", "# evaluate the model\n", "error =\n", "model.evaluate(X_test, y_test, verbose=0)\n", "print('MSE: %.3f, RMSE: %.3f % (error,\n", "sqrt(error)))\n", "# make a prediction\n", "row =\n", "[0.00632, 18.00, 2.310, 0, 0.5380, 6.5750, 65.20, 4.0900, 1, 296.0, 0.15, 30, 396.90, 4.98]\n", "yhat =\n", "model.predict([row])\n", "print('Predicted: %.3f % yhat')", "execution_count": null, "outputs": [{"output_type": "stream", "text": ["(339, 13) (167, 13) (339,) (167,)"], "name": "stdout"}]}, {"cell_type": "code", "metadata": {"colab": {"base_uri": "https://localhost:8080/"}, "id": "3hRyaZCryHri", "outputId": "1015f67d-bdc2-48ee-fa95-b5e3b90ef173"}, "source": ["#https://machinelearningmastery.com/how-to-stop-training-deep-neural-networks-at-the-right-time-using-early-stopping\n", "# mlp overfit on the moons dataset\n", "with simple early stopping\n", "from sklearn.datasets import make_moons\n", "from keras.models import Sequential\n", "from keras.layers import Dense\n", "from keras.callbacks import\n", "EarlyStopping\n", "from matplotlib import pyplot\n", "# generate 2d classification dataset\n", "X, y =\n", "make_moons(n_samples=100, noise=0.2, random_state=1)\n", "# split into train and\n", "test\n", "n_train = 30\n", "trainX, testX = X[:n_train, :], X[n_train:, :]\n", "trainy, testy = y[:n_train],\n", "y[n_train:]\n", "# define model\n", "model = Sequential()\n", "model.add(Dense(500, input_dim=2,\n", "activation='relu'))\n", "model.add(Dense(1,\n", "activation='sigmoid'))\n", "model.compile(loss='binary_crossentropy', optimizer='adam', metrics=[accuracy])\n", "# simple early stopping\n", "es = EarlyStopping(monitor='val_loss', mode='min',\n", "verbose=1)\n", "# fit model\n", "history = model.fit(trainX, trainy, validation_data=(testX, testy),\n", "epochs=4000, verbose=0, callbacks=[es])\n", "# evaluate the model\n", "_, train_acc =\n", "model.evaluate(trainX, trainy, verbose=0)\n", "_, test_acc = model.evaluate(testX, testy,\n", "verbose=0)\n", "print('Train: %.3f, Test: %.3f % (train_acc,\n", "test_acc))", "execution_count": null, "outputs": [{"output_type": "stream", "text": ["Epoch 00224: early\n", "stopping\n", "Train: 0.967, Test: 0.814\n"], "name": "stdout"}]}]
```